

Detecția anomaliilor în rețele folosind Machine Learning

1st Semester of 2023-2024

Author

`gabriel-lucian.tilica@s.unibuc.ro`

Abstract

This project focused on the detection of anomalies in network traffic using machine learning techniques. The key goal was to develop an AI-based system capable of identifying unusual or potentially malicious activities within network environments. The UNSW-NB15 dataset was employed to train and evaluate an Autoencoder model implemented in PyTorch, which identified anomalies based on reconstruction errors.

Throughout the project, significant learning outcomes were achieved, including data pre-processing, feature selection, and the practical implementation of neural network architectures. Notable methods such as Principal Component Analysis (PCA) and Recursive Feature Elimination (RFE) were used to enhance model efficiency and accuracy. The results highlighted the potential of Autoencoders for effective anomaly detection.

For future work, I aim to expand my knowledge by exploring Isolation Forests as an anomaly detection method and comparing them with Autoencoders to assess their strengths and weaknesses. Additionally, I plan to work with diverse datasets beyond UNSW-NB15 to evaluate the system's generalizability and enhance its adaptability to various network environments. Lastly, I intend to investigate ensemble learning techniques that combine multiple models to create more robust and scalable solutions for real-time anomaly detection.

1 Introduction

The problem I am trying to solve is the detection of anomalies in network traffic to identify unusual or potentially malicious activities. This issue is critical in cybersecurity and network management, as anomalies often indicate security breaches, malware activity, Distributed Denial of Service (DDoS) attacks, or misconfigurations within the network. The key aspects of the problem I am working at are the following. Modern networks generate massive amounts of traffic, which makes it challenging for traditional methods to effectively monitor and detect anomalies. At the same time, cyberattacks and malicious activities are becoming increasingly sophisticated, necessitating advanced detection mechanisms that go beyond static rule-based systems. Many network applications also require real-time or near-real-time anomaly detection to minimize potential damage. Moreover, existing systems often produce a high number of false positives, overwhelming administrators and making it difficult to focus on genuine threats. This problem is critically important because anomalies can serve as early indicators of security breaches or attacks. Promptly detecting and addressing network issues helps prevent downtime and improves overall network performance. Additionally, automated anomaly detection significantly reduces the need for manual monitoring, leading to considerable time and resource savings. By developing an AI-based system for anomaly detection, I aim to address these challenges using machine learning models that can learn patterns in network traffic and differentiate between normal and abnormal behavior. Since I'm the only contributor to this project, it was not easy, but I tried to approach this problem by presenting an implementation of code and documenting one machine learning model for anomaly detection: Autoencoder. Using the UNSW-NB15 dataset, I ensured that the model was reproducibly implemented, with clear instructions and comments in the code. The Autoencoder was used to detect anomalies by reconstructing input data and identifying deviations. Additionally, I created a detailed explanation and tutorial for the model, outlining its structure, how it process data, and his suit-

ability for anomaly detection. This includes insights gained from the results. The approach I am going to take involves a thorough phase of research and study before starting the implementation. During this phase, I dedicated time to documenting and planning the project. This included learning about network traffic flows and their characteristics, such as protocols (e.g., TCP/IP, HTTP, DNS) and essential parameters like IP addresses, ports, latency, and packet size. Understanding these elements was crucial for interpreting the dataset and identifying patterns in network behavior. I focused on studying the principles behind anomaly detection models, particularly Autoencoders for reconstructing input patterns. Additionally, I explored related methods, such as clustering techniques and recurrent neural networks (RNNs), to understand their potential applicability in this domain. While I also reviewed Isolation Forests conceptually, they were not part of the implementation in this project. Furthermore, I analyzed publicly available datasets like UNSW-NB15 to familiarize myself with their structure and relevance to the problem. This included identifying features critical for anomaly detection and understanding how these features align with the behavior of the Autoencoder model I implemented. In the implementation phase, I focused on developing an Autoencoder for detecting anomalies in network traffic. First, I preprocessed the UNSW-NB15 dataset by extracting relevant features and normalizing the data to ensure compatibility with the model. The Autoencoder was trained to reconstruct normal network traffic patterns, with anomalies identified as inputs that deviated significantly from the reconstruction. The implementation was done in Python, with detailed documentation and clear instructions provided to ensure reproducibility. I evaluated the performance of the Autoencoder using metrics such as precision, recall, and F1-score, and generated visualizations to illustrate its effectiveness in detecting anomalies. Lastly, I compiled the results into a concise tutorial and presentation, highlighting the strengths and limitations of the Autoencoder approach. I chose to approach this project because I have a strong passion for cybersecurity and plan

135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185

to further my studies in this field by pursuing a master's degree in Security and Applied Logic. This project aligns perfectly with my interests, allowing me to explore practical applications of machine learning in enhancing network security. Cybersecurity is a critical domain where the complexity and volume of network traffic make traditional monitoring methods less effective. The integration of AI in this field is transformative, as it enables systems to adapt to evolving threats, identify anomalies in real-time, and reduce false positives through intelligent pattern recognition. By leveraging AI techniques like Autoencoders and Isolation Forests, this project demonstrates how machine learning can provide advanced, scalable solutions for detecting malicious activities in networks. This approach not only deepens my understanding of AI's role in cybersecurity but also serves as a foundation for further academic and professional growth in this domain. Throughout this project, I gained a deeper understanding of machine learning techniques and their applications in anomaly detection. I learned how to preprocess and analyze network traffic data, particularly using the UNSW-NB15 dataset. This involved understanding the structure of network traffic, selecting relevant features, and applying normalization and dimensionality reduction techniques to prepare the data effectively. Implementing an Autoencoder model in PyTorch allowed me to detect anomalies based on reconstruction errors, which further enhanced my understanding of neural network architectures, training processes, and evaluation metrics. I also developed a strong appreciation for the importance of feature selection and its impact on model performance. By applying techniques like Principal Component Analysis (PCA) and Recursive Feature Elimination (RFE), I was able to improve the system's efficiency and accuracy. Additionally, I gained valuable experience in visualizing data and model performance using techniques like kernel density plots and loss curves, which helped me interpret the results and refine the anomaly detection system. In the future, I would like to expand my knowledge by studying Isolation Forests, a powerful anomaly detection method that iso-

lates data points in feature space to identify outliers. Learning about this technique will allow me to compare it with Autoencoders and assess its suitability for different scenarios. I also plan to explore diverse datasets beyond UNSW-NB15 to evaluate the generalizability of my system and adapt it to different network environments. Furthermore, I aim to delve into ensemble learning techniques, such as combining Isolation Forests with neural networks or other models, to create more accurate and robust systems for real-time anomaly detection. By reflecting on my experiences in this project and setting clear goals for future learning, I hope to continue improving my understanding and application of machine learning techniques in the field of cybersecurity.	186 187 188 189 190 191 192 193 194 195 196 197 198 199 200 201 202 203
2 Approach	204
2.1 Learning Phase	205
Before starting the implementation, I dedicated time to understanding the foundational concepts and tools required for this project. This process was crucial to ensure that I could effectively develop and evaluate the anomaly detection system. Below, I outline the key aspects of this learning phase:	206 207 208 209 210 211
2.1.1 Understanding Network Traffic	212
I began by studying the characteristics of network traffic to understand the nature of the data I would be working with. This included:	213 214 215
<ul style="list-style-type: none">• Learning about network protocols such as TCP/IP, HTTP, and DNS, which govern communication in networks.• Understanding essential parameters of network traffic, including IP addresses, ports, packet size, and latency. These features play a critical role in identifying patterns and detecting anomalies.	216 217 218 219 220 221 222 223
2.1.2 Exploring Machine Learning for Anomaly Detection	224 225
Next, I focused on learning about machine learning models commonly used for anomaly detection:	226 227
<ul style="list-style-type: none">• Autoencoders: I studied how autoencoders reconstruct input data and how deviations from reconstruction indicate anomalies. This included understanding the encoder-decoder architecture and its training process. Autoen-	228 229 230 231 232

233
234
235
236
237
238

239
240
241
242
243
244

245
246
247
248

249
250

251
252

253
254
255
256
257
258
259
260

261

262
263
264
265
266

267
268

269
270
271
272
273
274
275
276
277
278
279

coders is a type of neural network used in unsupervised learning to encode input data into a compressed representation and then reconstruct it. Anomalies are detected as inputs that deviate significantly from the reconstructed data, indicating unusual patterns.

- **Other Techniques:** Although not implemented in this project, I also explored related methods such as Isolation Forests, clustering techniques, and Recurrent Neural Networks (RNNs) to understand their strengths and limitations.

2.1.3 Familiarizing with Public Datasets

To ensure the system was trained on relevant and diverse data, I analyzed publicly available datasets like UNSW-NB15. This involved:

- Understanding the structure and features of the dataset.
- Identifying which features were most relevant for detecting anomalies in network traffic.

The UNSW-NB15 dataset is a comprehensive collection of network traffic data designed to evaluate intrusion detection systems. It encompasses both normal activities and a variety of contemporary synthesized attack behaviors, making it suitable for developing and testing machine learning models for anomaly detection.

Key Features of the UNSW-NB15 Dataset:

Data Composition: The dataset includes a total of 2,540,044 records, divided into four CSV files. Specifically, the training set comprises 175,341 records, while the testing set contains 82,332 records.

Attack Categories: It features nine distinct types of attacks:

Fuzzers: Techniques that send unexpected or random input data to applications to discover vulnerabilities. **Analysis:** Methods involving the gathering and studying of information to exploit system vulnerabilities. **Backdoors:** Unauthorized access points created by inserting malicious code, allowing attackers to bypass security measures. **DoS (Denial of Service):** Attacks aimed at making a machine or network resource unavailable to its intended users. **Exploits:** Attacks that take advantage

of software vulnerabilities to gain unauthorized access. **Generic:** Attacks that can be applied to various platforms without modification. **Reconnaissance:** Techniques used to gather information about a system to find ways to infiltrate it. **Shellcode:** Small pieces of code used as the payload in the exploitation of a software vulnerability. **Worms:** Malicious software programs that replicate themselves to spread to other computers. **Feature Set:** The dataset contains 49 features extracted using Argus and Bro-IDS tools, which are categorized into:

Flow Features: Attributes related to the flow of data packets between source and destination. **Basic Features:** Fundamental attributes such as protocol type, service, and duration. **Content Features:** Information derived from the data portion of the packets, like the number of failed login attempts. **Time Features:** Attributes related to the timing of the connections, such as the time to live (TTL) of the packets. **Additional Generated Features:** Features created to provide more insights, like the number of compromised conditions.

2.2 Implementation Phase

After completing the learning phase, I proceeded to implement the project by setting up the necessary environment and adding core files. The following steps outline the implementation process:

2.2.1 Project Setup

The project was developed in Visual Studio Code (VSCode), providing a structured development environment to manage, edit, and execute the code. The initial setup focused on preparing the environment and incorporating essential files for anomaly detection.

- **Adding the Autoencoder Implementation:** The primary file added to the project is an existing implementation of an Autoencoder-based anomaly detection system tailored to the UNSW-NB15 dataset. This file, sourced from https://github.com/alik604/cyber-security/blob/master/Intrusion-Detection/UNSW_NB15%20-%20PyTorch%20MLP%20and%20autoEncoder.ipynb, contains:

280
281
282
283
284
285
286
287
288
289
290
291
292

293
294
295
296
297
298
299
300
301
302
303
304

305
306
307
308
309

310
311
312
313
314
315
316

317
318
319
320
321
322
323
324
325
326
327
328

329	– Data preprocessing steps specific to the	These steps ensure that the project en-	374
330	UNSW-NB15 dataset, including feature	vironment can be replicated on any ma-	375
331	extraction and normalization.	chine.	376
332	– An Autoencoder architecture imple-		
333	mented in PyTorch, designed to recon-	2.3 Understanding the Provided ipynb	377
334	struct network traffic patterns and detect	file	378
335	anomalies based on reconstruction loss.		
336	– Visualization and evaluation metrics to	After setting up the project environment,	379
337	assess the model’s effectiveness in iden-	the next step was to incorporate and under-	380
338	tifying anomalies.	stand the implementation provided in the	381
339		notebook UNSW_NB15 - PyTorch MLP	382
340	This file was integrated into the project as the	and autoEncoder.ipynb. This file	383
	core implementation for detecting anomalies.	contains the core implementation for anomaly	384
		detection using an Autoencoder. The follow-	385
341	• Creating the requirements.txt File:	ing steps were undertaken to fully compre-	386
342	To manage dependencies and ensure repro-	hend and adapt this notebook for the project:	387
343	ducibility, a requirements.txt file was		
344	created. This file lists all the Python libraries	2.3.1 Setup and Library Imports	388
345	necessary for running the project, including:		
346	– torch for building and training the Au-	The provided code begins by configuring	389
347	toencoder.	the environment and importing the neces-	390
348	– scikit-learn for preprocessing and	sary libraries for developing an anomaly	391
349	evaluation metrics.	detection system. Each section of this	392
350	– tabulate for presenting results in a	setup plays a specific role in ensuring	393
351	structured format.	the proper execution of the project. First,	394
352	– pandas and numpy for handling and	the following line is used to enhance the	395
353	manipulating the dataset.	Jupyter Notebook environment:	396
354	– matplotlib for visualizations.	%config IPCompleter.greedy=True	397
355	– ipykernel to enable running the note-	This configuration enables comprehen-	398
356	book in a Jupyter environment.	sive auto-completion features for meth-	399
357		ods, attributes, and variables, streamlin-	400
358	This file ensures that all dependencies can be	ing the development process and reduc-	401
	installed with a single command.	ing potential errors.	402
		Next, several essential Python libraries	403
359	• Creating the setup.txt File: A	are imported:	404
360	setup.txt file was created to provide clear	* Pandas: Facilitates the manipulation	405
361	and concise instructions for setting up and	and analysis of structured data in tab-	406
362	running the project. The file includes:	ular form.	407
363	– Instructions to create a virtual environ-	* NumPy: Provides tools for efficient	408
364	ment using:	numerical computations and array	409
365	python3 -m venv myenv	operations.	410
366		* SciPy: Adds advanced mathematical	411
367	– Steps to activate the virtual environment:	functions useful for scientific com-	412
368	source myenv/bin/activate	puting.	413
369		* Matplotlib: Includes tools for	414
370	– Commands to install the required li-	creating data visualizations. The	415
371	braries:	%matplotlib inline com-	416
372	pip3 install -r requirements.txt	mand ensures that all plots are	417
373		displayed directly within the	418
		notebook.	419
		* Scikit-learn: Offers a suite	420
		of machine learning tools for	421

Identifying Low-Variance and Low-Correlation Features

This code segment focuses on identifying features in the dataset with minimal variance and low correlation to the target variable, 'attackcat'.

First, a copy of the 'combined_data' DataFrame is created to ensure the original dataset remains unchanged during processing.

The low-variance features are identified using the standard deviation of each column. The 'std()' method computes the standard deviation for all features, and the 'nsmallest()' function selects the seven features with the smallest values. These features are stored in the 'lowSTD' list.

Next, the code calculates the absolute correlation of each feature with the 'attackcat' column using the 'corr()' method. Features with the lowest correlation to 'attackcat' are identified using the 'nsmallest()' function and stored in the 'lowCORR' list. This step helps pinpoint features that are least relevant to predicting the target variable.

The identified low-variance and low-correlation features may be considered for removal or further analysis to improve model performance by reducing noise and redundancy.

Performing Dimensionality Reduction and Updating the Dataset

This section of the code is responsible for performing dimensionality reduction on selected features in the dataset and subsequently updating the dataset with reduced features. The key steps are as follows: Initially, a list named `exclude` is created by combining features with low correlation (`lowCORR`) and low standard deviation (`lowSTD`). The `attack_cat` column, which serves as the target variable, is explicitly removed from this list if it exists. This ensures that the target variable is not included in the dimensionality reduction process.

The code then prints the shape of the dataset before dimensionality reduction and displays the features identified for replacement with Principal Component Analysis (PCA).

Next, PCA is applied to the columns specified in the `exclude` list, with the number of principal components set to three. This process reduces the dimensionality of the dataset while retaining as much variance as possible. The `explained_variance_ratio_` is calculated and printed, providing insight into how much variance is captured by the reduced features.

After PCA is performed, the features listed in `exclude` are dropped from the dataset using the `drop` method. The reduced features are stored in a DataFrame and joined with the updated dataset to maintain a consistent structure.

Finally, the code prints the shape of the dataset after dimensionality reduction, showcasing the changes made to the dataset. The updated dataset now includes the reduced features in place of the original ones, which simplifies the dataset while preserving its predictive power.

Scaling the Duration Feature

This section of the code performs a scaling operation on the `dur` (duration) feature within the dataset. The primary objective is to transform the values of this feature to a larger scale for better numerical representation during model training and analysis.

The operation scales the `dur` feature by multiplying its values by 10,000. This scaling ensures that the feature has a more significant numerical range, which might help certain machine learning models that are sensitive to the magnitude of feature values.

After scaling, the code prints a message indicating that the `dur` feature has been scaled up by a factor of 10,000. Additionally, the `head()` method is called on the `combined_data` DataFrame to display the first few rows of the dataset, allowing verification of the scaling operation.

Preparing the Data for Model Training

This section of the code prepares the dataset for model training by splitting it into features (`data_x`) and labels (`data_y`), followed by normalizing the feature values.

The process begins by printing the shape of the original `combined_data` DataFrame. The `data_x` variable is created by dropping the `attack_cat` column, which serves as the label for classification. The `data_y` variable is assigned the `attack_cat` column, representing the target variable for the machine learning model. The shapes of both `data_x` and `data_y` are printed to verify the separation.

Normalization is applied to the `data_x` features using a lambda function. This function scales each feature to the range [0, 1] by subtracting the minimum value and dividing by the range (maximum - minimum). This ensures that all features are on

a consistent scale, which is essential for machine learning models sensitive to feature magnitudes. Two commented-out lines provide alternative normalization techniques:

```
* MinMaxScaler(): Normal-
    izes data to the range [0, 1]
    and is noted as better suited for
    VotingClassifier.
* StandardScaler(): Standard-
    izes data to have a mean of 0 and
    a standard deviation of 1.
```

These options provide flexibility depending on the requirements of different machine learning models or algorithms.

Splitting the Data into Training and Testing Sets

This portion of the code splits the prepared dataset into training and testing subsets using the `train_test_split` function from the `sklearn.model_selection` module. The variables `X_train` and `X_test` represent the feature data for training and testing, while `y_train` and `y_test` represent the corresponding labels. The split is configured with the following parameters:

```
* test_size=0.50: Specifies that
    50% of the data will be allocated to
    the test set, and the remaining 50%
    will be used for training. The com-
    ment suggests that the test_size
    value can be adjusted depending on
    computational constraints and how
    long the user is willing to wait for
    processing.
* random_state=42: Ensures re-
    producibility of the split by using a
    fixed random seed.
```

This step is crucial for evaluating the model's performance, as it creates separate datasets for training the model and testing its accuracy on unseen data. The split ensures that the evaluation is unbiased and reflects the model's generalization capabilities.

Implementing and Evaluating Ensemble and Individual Models

This portion of the code implements three individual classifiers and an ensemble model to evaluate their performance on the training and testing datasets:

1. Classifiers Used:

```
* DecisionTreeClassifier
    (DTC): A simple decision tree-
    based model that splits the data
    recursively based on features to
    make predictions.
* RandomForestClassifier
    (RFC): An ensemble method that
    combines multiple decision trees,
    with 50 trees specified via the
    n_estimators parameter and
    a random seed for reproducibility
    (random_state=1).
* ExtraTreesClassifier
    (ETC): Another ensem-
    ble method similar to
    RandomForestClassifier,
    with 75 trees
    (n_estimators=75) and
    customized hyperparameters, such
    as the gini criterion for split-
    ting and no bootstrap sampling
    (bootstrap=False).
```

2. Voting Classifier: The ensemble model, `VotingClassifier`, combines the predictions of the three classifiers mentioned above. It uses `voting='hard'`, which means the final prediction is based on majority voting among the classifiers.

3. Training and Evaluation: The `for` loop iterates through the three individual classifiers (DTC, RFC, and ETC) and the ensemble model (`ecclf`). For each classifier:

```
* clf.fit(X_train,
    y_train): Trains the classi-
    fier on the training data.
* clf.score(X_test,
    y_test): Evaluates the clas-
    sifier on the test data and returns the
    accuracy score.
* The accuracy score is printed with
    the classifier's name for comparison:
    Acc: %0.7f [%s].
```

This step allows for a direct comparison of the performance of individual classifiers and the ensemble model, providing insights into which approach performs best for the given dataset.

Feature Selection and Model Evaluation

This code combines feature selection techniques with ensemble and individual model evaluations to improve classification performance.

1. Feature Selection with Recursive Feature Elimination (RFE): The RFE (Recursive Feature Elimination) method is used to select the top 10 most important features from the dataset:

```
* RFE(DecisionTreeClassifier(
10).fit(X_train,
y_train): This initial-
izes an RFE object with a
DecisionTreeClassifier
as the estimator and identifies
the 10 most significant features
by iteratively removing the least
important ones.
* np.where(rfe.support_==True)
Identifies the indices of the selected
features.
* The selected feature names are stored
in the variable whitelist, and
the training and testing datasets
are reduced to these features us-
ing X_train[whitelist] and
X_test[whitelist].
```

2. Model Initialization: Three classifiers are defined:

```
* DecisionTreeClassifier
(DTC): A simple decision tree
model.
* RandomForestClassifier
(RFC): An ensemble of 50 deci-
sion trees with a random seed for
reproducibility.
* ExtraTreesClassifier
(ETC): An ensemble method with
75 decision trees, using gini as the
splitting criterion and no bootstrap
sampling.
```

3. Evaluation with Original and Reduced Features:

```
* The classifiers are trained and
evaluated twice: once with the
full feature set (X_train and
X_test) and once with the re-
duced feature set (X_train_RFE
and X_test_RFE).
* The accuracy scores for the classi-
fiers are compared to evaluate the im-
pact of feature selection on perfor-
mance.
```

4. Voting Classifier: The ensemble model VotingClassifier combines the predictions of the three classifiers using majority voting

(voting='hard'). The performance of the ensemble model is also evaluated.

5. Output: For each classifier (including the ensemble), the accuracy score is printed in the format `Acc: %0.7f [%s]` to compare the effectiveness of the classifiers on the full and reduced feature sets.

This code demonstrates the integration of feature selection techniques with machine learning models to enhance classification accuracy and reduce model complexity.

Inspecting Feature-Reduced Dataset and Target Labels

This code performs two actions to inspect the data after applying feature selection and to understand the distribution of target labels:

1. Checking the Shape of the Feature-Reduced Training Dataset:

```
* X_train_RFE.shape prints the
dimensions of the feature-reduced
training dataset after applying Re-
cursive Feature Elimination (RFE).
This provides insight into how many
features and samples are retained for
model training.
```

2. Inspecting the Unique Labels in the Target Variable:

```
* set(y_train) returns the
unique values in the target variable
y_train. This reveals the distinct
classes present in the training labels,
ensuring that all expected classes are
accounted for.
```

These actions are essential for verifying the data preparation process and confirming that the feature reduction and label assignments are correctly handled.

Defining the Neural Network Architecture

This segment of code defines a fully connected neural network (FCNN) for anomaly detection, specifying its architecture, parameters, and behavior.

1. Selecting the Device: The code checks if a CUDA-compatible GPU is available for computation. If a GPU is available, the device is set to 'cuda'; otherwise, it defaults to 'cpu'. The second assignment, `device = 'cpu'`, explicitly overrides this, ensuring that the model runs on the CPU.

2. Setting Hyperparameters: The code initializes several hyperparameters critical to training the neural network:

- * `input_size`: The number of input features, set to 10.
- * `hidden_size` and `hidden_size_2`: The sizes of the two hidden layers, each with 64 neurons.
- * `num_classes`: The number of output classes, set to 10.
- * `num_epochs`: The number of training iterations, set to 40.
- * `batch_size`: The size of mini-batches used during training, set to 32.
- * `learning_rate`: The step size for the optimizer, set to 0.001.

3. Defining the Neural Network Class: The `NeuralNet` class defines a custom neural network architecture using the PyTorch `nn.Module` base class:

- * The `__init__()` method initializes the layers and activation functions:
 - `self.l1`: A fully connected layer (`nn.Linear`) connecting the input layer to the first hidden layer.
 - `self.l2`: A fully connected layer connecting the first hidden layer to the second hidden layer.
 - `self.l3`: A fully connected layer connecting the second hidden layer to the output layer.
 - `self.relu`: A Rectified Linear Unit (ReLU) activation function applied after the first and second layers to introduce non-linearity.
 - `self.elu`: An Exponential Linear Unit (ELU) activation function, though it is defined but not currently used in the `forward()` method.
- * The `forward()` method defines the forward pass of the network:
 - The input passes through the first layer (l1) and then through the ReLU activation function.
 - The output of the first layer is

passed through the second layer (l2) and another ReLU activation. Finally, the data is passed through the third layer (l3) to produce the network's output. Note that no activation function or softmax is applied at the output, as it is likely handled separately, depending on the task.

This modular architecture facilitates easy experimentation with different activation functions, layer sizes, and other hyperparameters.

Training the Neural Network

This segment of the code is responsible for training the neural network model on the reduced feature set using the selected hyperparameters, loss function, and optimizer.

1. Model Initialization:

- * The `NeuralNet` model is instantiated using the previously defined architecture and configured with the `input_size`, `hidden_size`, and `num_classes` parameters. The model is moved to the specified device (CPU in this case).

2. Loss Function:

- * The loss function used is `nn.CrossEntropyLoss`, which combines `nn.LogSoftmax` and `nn.NLLLoss` in one class. It is well-suited for multi-class classification problems.

3. Optimizer:

- * The `torch.optim.Adam` optimizer is selected to update the model parameters. It uses an adaptive learning rate (`learning_rate=0.001`) to optimize the training process efficiently.

4. Training Loop:

- * The training process runs for a specified number of epochs (`num_epochs=40`).
- * The training data is iterated in mini-batches, with the batch size set to `batch_size=32`.
- * For each mini-batch:
 - The input features (`x`) and corresponding labels (`y`) are converted

1011	· self.l1: A linear layer mapping the input size to the hidden_size.	Decoder attempts to reconstruct the original input from this compressed representation.	1062
1012			1063
1013			1064
1014	· self.l2: A linear layer mapping the hidden_size to the center_size, which represents the bottleneck or compressed feature space.	* These classes are essential for building and training an autoencoder, a type of neural network used for unsupervised learning and dimensionality reduction.	1065
1015			1066
1016			1067
1017			1068
1018			1069
1019	· self.acti: An Exponential Linear Unit (ELU) activation function introducing non-linearity.		
1020		Defining the Recurrent Autoencoder Class	1070
1021	* forward(): Defines the forward pass of the encoder:	The RecurrentAutoencoder class combines the previously defined Encoder and Decoder classes into a complete autoencoder architecture. It is built using the PyTorch nn.Module base class.	1071
1022			1072
1023	· Input x is passed through the first linear layer (l1).		1073
1024		1. Class Initialization:	1074
1025			1075
1026	· The activation function (acti) is applied to introduce non-linearity.	* The __init__() method initializes the encoder and decoder components of the autoencoder:	1076
1027			1077
1028	· The transformed data is passed through the second linear layer (l2), producing the compressed representation.	· self.encoder: An instance of the Encoder class, which compresses the input data into a lower-dimensional representation.	1078
1029			1079
1030		· self.decoder: An instance of the Decoder class, which reconstructs the original input from the compressed representation.	1080
1031			1081
1032	3. Decoder Class: The Decoder class reconstructs the original data from the compressed representation.	· Both the encoder and decoder are moved to the specified device (e.g., 'cpu').	1082
1033			1083
1034			1084
1035	* __init__(): Initializes two fully connected layers and an activation function:		1085
1036			1086
1037			1087
1038	· self.l1: A linear layer mapping the center_size back to the hidden_size.		1088
1039			1089
1040	· self.l2: A linear layer mapping the hidden_size to the output_size, which matches the original input dimension.	* The architecture of the encoder and decoder is printed for verification using print(self.encoder) and print(self.decoder).	1090
1041			1091
1042	· self.acti: An Exponential Linear Unit (ELU) activation function for non-linearity.		1092
1043			1093
1044		2. Forward Pass:	1094
1045			1095
1046	* forward(): Defines the forward pass of the decoder:	* The forward() method defines the forward propagation logic:	1096
1047			1097
1048	· Input x (compressed representation) is passed through the first linear layer (l1).	· The input data (x) is passed through the encoder to obtain a compressed representation (encoded).	1098
1049			1099
1050	· The activation function (acti) is applied.	· The compressed representation is then passed through the decoder to produce the reconstructed data (decoded).	1100
1051			1101
1052	· The transformed data is passed through the second linear layer (l2) to reconstruct the original input.	· The output of the decoder (decoded) is returned as the final result.	1102
1053			1103
1054			1104
1055			1105
1056			1106
1057			1107
1058		3. Summary:	1108
1059	4. Summary:		1109
1060	* The Encoder reduces the dimensionality of the input data, while the	* The RecurrentAutoencoder encapsulates the complete autoencoder structure by combining an	1110
1061			1111

1112	encoder and a decoder. It processes		
1113	input data, compresses it into a latent		
1114	representation, and reconstructs the		
1115	original input.		
1116	* This class serves as the core archi-		
1117	tecture for training and testing the		
1118	autoencoder on a given dataset.		
1119	Training the Autoencoder Model		
1120	The <code>train_model</code> function is responsible for		
1121	training the autoencoder model using the provided		
1122	training and validation datasets. This function in-		
1123	cludes mechanisms for optimization, loss evalua-		
1124	tion, and tracking performance across epochs.		
1125	Function Parameters:		
1126	* <code>model</code> : The autoencoder model to		
1127	be trained.		
1128	* <code>train_dataset</code> : The dataset		
1129	used for training the model.		
1130	* <code>val_dataset</code> : The dataset used		
1131	for validating the model's perfor-		
1132	mance.		
1133	* <code>n_epochs</code> : The number of epochs		
1134	for training.		
1135	* <code>batch_size</code> : The size of each		
1136	batch used during training and val-		
1137	idation.		
1138	* <code>lr</code> : The learning rate for the opti-		
1139	mizer.		
1140	Training Procedure:		
1141	1. Initialization:		
1142	* An Adam optimizer is initialized		
1143	with the model parameters and the		
1144	specified learning rate (<code>lr</code>).		
1145	* A learning rate scheduler		
1146	(<code>ReduceLROnPlateau</code>) ad-		
1147	justs the learning rate based on		
1148	validation loss, reducing it if no		
1149	improvement is observed for 3		
1150	epochs.		
1151	* The loss function used is Mean		
1152	Squared Error (<code>MSELoss</code>) with		
1153	reduction set to <code>sum</code> .		
1154	* <code>history</code> : A dictionary is cre-		
1155	ated to track training and valida-		
1156	tion losses across epochs.		
1157	* <code>best_model_wts</code> : A copy		
1158	of the model's initial weights		
1159	is stored to retain the best-		
1160	performing model.		
	2. Epoch Loop: The training process		1161
	iterates for the specified number of		1162
	epochs (<code>n_epochs</code>):		1163
	* Training Phase:		1164
	· The model is set to training mode		1165
	using <code>model.train()</code> .		1166
	· For each batch in the training		1167
	dataset:		1168
	· The true sequence (<code>seq_true</code>)		1169
	is fed into the model to generate		1170
	predictions (<code>seq_pred</code>).		1171
	· The loss between predicted and		1172
	true sequences is calculated using		1173
	<code>criterion</code> .		1174
	· Gradients are computed via back-		1175
	propagation, and the optimizer		1176
	updates the model's weights.		1177
	· The training loss for each batch		1178
	is recorded and averaged for the		1179
	epoch.		1180
	* Validation Phase:		1181
	· The model is set to evaluation		1182
	mode using <code>model.eval()</code> .		1183
	· For each batch in the validation		1184
	dataset:		1185
	· The true sequence is passed		1186
	through the model to generate		1187
	predictions.		1188
	· The loss is computed and		1189
	recorded for each batch.		1190
	· The validation loss for each		1191
	batch is averaged for the epoch.		1192
	3. Loss Tracking and Scheduler Up-		1193
	date:		1194
	* Training and validation losses		1195
	are stored in the <code>history</code> dic-		1196
	tionary.		1197
	* The learning rate scheduler ad-		1198
	justs the learning rate based on		1199
	the validation loss.		1200
	* If the current validation loss is		1201
	lower than the best observed		1202
	loss, the model's weights are		1203
	saved as the best weights.		1204
	4. Model Finalization: After all		1205
	epochs, the model's weights are re-		1206
	stored to the best-performing config-		1207
	uration, ensuring the returned model		1208
	represents the best version encoun-		1209
	tered during training.		1210

1211
1212
1213
1214
1215

1216
1217
1218
1219
1220
1221
1222
1223
1224
1225

1226
1227
1228

1229
1230
1231
1232
1233

1234
1235
1236
1237

1238

1239
1240
1241
1242
1243
1244

1245
1246

1247
1248
1249

1250
1251
1252

Output:

- * The function returns the trained model in evaluation mode and the `history` dictionary containing training and validation loss trends.

2.4 Splitting and Normalizing the Dataset

In this section of the code, the dataset is split into two subsets: one representing *normal* network traffic and the other representing *attack* traffic. Additionally, normalization is applied to ensure that features are scaled for consistency in subsequent analysis. An optional experimental feature selection process is included, controlled by the `EXPERIMENT` flag.

Overview of the Code:

- * The `EXPERIMENT` variable determines whether to apply additional filtering to the dataset.
- * If `EXPERIMENT = True`, a statistical analysis of feature differences between *normal* and *attack* traffic is conducted to retain only significant features.
- * If `EXPERIMENT = False`, the dataset is simply split into two subsets, and min-max normalization is applied.

Steps When `EXPERIMENT = True`:

1. The dataset is divided into two groups:
 - * *Normal traffic*: Rows where `attack_cat` equals 6.
 - * *Attack traffic*: Rows where `attack_cat` is not 6.
2. Both groups are normalized using min-max scaling.
3. Statistical summaries (mean and `std`) are computed for each feature in the two subsets.
4. The absolute differences in mean and standard deviation between the two subsets are calculated.

5. Features with a mean difference greater than 0.1 are retained. These features are then used to filter the *normal* and *attack* datasets.

Steps When `EXPERIMENT = False`:

1. The dataset is split into *normal* and *attack* subsets based on the value of `attack_cat`.
2. Min-max normalization is applied to both subsets, scaling all features to a range of 0 to 1.

Final Output:

- * The variable `normal` contains rows corresponding to normal network traffic with the `attack_cat` column removed.
- * The variable `attack` contains rows corresponding to attack traffic with the `attack_cat` column removed.
- * The dimensions of the `normal` subset are printed for verification.

This step ensures that the dataset is prepared for further analysis or modeling by segregating the data and applying normalization. Additionally, the optional feature selection process helps reduce dimensionality and retain only the most relevant features for anomaly detection.

2.5 Preparing Data for Training and Testing

In this section, the dataset is further divided and converted into tensors for use in a PyTorch-based model. The code processes the *normal* and *attack* subsets to create training and testing datasets specifically for normal traffic.

Steps in the Code:

1. The `normal` and `attack` datasets, initially stored as pandas DataFrames, are converted to NumPy arrays using `.values`. This conversion is required to prepare the data for PyTorch tensor operations.
2. The `normal` dataset is split into two parts:

1253
1254
1255
1256

1257

1258
1259
1260

1261
1262
1263

1264

1265
1266
1267
1268

1269
1270
1271
1272

1273
1274

1275
1276
1277
1278
1279
1280

1281
1282
1283
1284
1285
1286
1287

1288
1289
1290
1291
1292
1293
1294

1295
1296

1297	* normal_train: 90% of the	ensures that the input and output	1343
1298	data is allocated for training.	layers of the autoencoder match the	1344
1299	* normal_test: The remain-	dimensionality of the data.	1345
1300	ing 10% is reserved for testing.		
1301	3. The training and testing subsets are	2. Hidden and Center Sizes:	1346
1302	converted to PyTorch tensors using	* hidden_size is set to 7,	1347
1303	torch.as_tensor(), with the	defining the size of the hidden	1348
1304	data type specified as float for	layers in the encoder and de-	1349
1305	numerical computations. These ten-	coder.	1350
1306	sors are moved to the device (CPU	* center_size is set to 2,	1351
1307	or GPU) specified earlier in the	representing the size of the la-	1352
1308	script.	tent representation learned by	1353
1309	4. Additionally, the entire normal	the autoencoder.	1354
1310	and attack datasets are	3. Training Hyperparameters:	1355
1311	converted into PyTorch ten-	* num_epochs: The number	1356
1312	sors (normal_tensor and	of epochs for training is set to	1357
1313	attack_tensor), which may	25.	1358
1314	be used for evaluation or analysis	* batch_size: The batch	1359
1315	purposes.	size for training is set to 16,	1360
1316		defining how many samples	1361
	Output:	are processed together in a	1362
1317	* The shapes of attack_tensor,	single forward and backward	1363
1318	normal, x_train, and x_test	pass.	1364
1319	are printed to verify the dimensions	* lr: The learning rate is set to	1365
1320	of the data at each stage of process-	0.01, controlling the step size	1366
1321	ing.	of the optimizer.	1367
1322	* This ensures that the datasets are	4. Model Initialization: A	1368
1323	correctly divided and converted into	RecurrentAutoencoder	1369
1324	tensors, which are compatible with	instance is created using the defined	1370
1325	PyTorch models.	parameters for input_size,	1371
1326		output_size, hidden_size,	1372
1327	This step prepares the data for subsequent train-	and center_size. The model	1373
1328	ing and testing phases, ensuring that the model	is moved to the specified device	1374
1329	receives appropriately formatted input for anomaly	(CPU or GPU).	1375
	detection.		
1330	2.6 Training the Recurrent Autoencoder	5. Training the Model: The	1376
1331	Model	train_model function is called	1377
1332		with the following arguments:	1378
1333	This section describes the setup and training of the	* model: The initialized Recur-	1379
1334	Recurrent Autoencoder model. The code initializes	rent Autoencoder.	1380
1335	model parameters, trains the model, and saves it	* x_train and x_test:	1381
1336	for future use.	Training and testing datasets	1382
	Steps in the Code:	prepared earlier.	1383
1337	1. Input and Output Dimen-	* n_epochs, batch_size,	1384
1338	sions: The input_size and	and lr: The defined training	1385
1339	output_size are both set to	hyperparameters.	1386
1340	the number of features in the		
1341	normal dataset, obtained using	This function trains the model using	1387
1342	normal.shape[1]. This	Mean Squared Error (MSE) as the	1388

1389	loss function and returns the trained	* It helps identify if the model is converging by observing whether the	1434
1390	model along with a history of training	loss values decrease over epochs.	1435
1391	and validation losses for analysis.		1436
1392			
1393	6. Saving the Model: After training,	* A significant gap between training	1437
1394	the model is saved to a file named	and validation loss may indicate	1438
1395	<code>model.pth</code> using <code>torch.save</code> .	overfitting.	1439
1396	This allows the trained model to be	* Similar trends in training and validation	1440
1397	loaded and reused without retraining.	loss suggest the model generalizes well to unseen data.	1441
1398			1442
1399	Purpose of this Step: This step trains the Recur-	2.8 Prediction and Loss Calculation	1443
1400	rent Autoencoder to learn latent representations of	The following code defines a function named	1444
1401	the normal network traffic. The trained model	<code>predict</code> , which is responsible for generating predictions	1445
1402	will be used in subsequent phases to detect anomalies	and calculating the corresponding reconstruction loss for a given dataset. This function is	1446
1403	based on deviations from normal patterns.	typically used during the evaluation phase to assess	1447
1404		the model's performance.	1448
1405	2.7 Visualizing Training and Validation Loss	Description of the Code:	1449
1406	This code snippet generates a visualization of the		1450
1407	training and validation loss over the course of the	* The function <code>predict</code> takes three	1451
1408	training epochs. The plot helps evaluate how well	arguments:	1452
1409	the model has learned during training and whether		
1410	overfitting or underfitting has occurred.	· <code>model</code> : The trained autoencoder model used for predictions.	1453
1411	Description of the Code:		1454
1412	* A figure and its axes are created	· <code>batch_size</code> : The number of samples processed in a single forward pass, which helps	1455
1413	using <code>plt.figure().gca()</code> , where <code>gca</code> returns the current axes.	manage memory usage.	1456
1414	* The training loss, stored in <code>history['train']</code> , is plotted	· <code>dataset</code> : The input data for which predictions and losses	1457
1415	over the epochs.	are computed.	1458
1416			1459
1417	* The validation loss, stored in <code>history['val']</code> , is plotted for	* Two lists, <code>predictions</code> and <code>losses</code> , are initialized to store the	1460
1418	comparison.	predicted outputs and corresponding reconstruction losses for each	1461
1419		batch.	1462
1420	* Labels for the x-axis (Epoch) and y-axis (Loss) are added to describe		
1421	the plot.	* The loss function is defined	1463
1422		using <code>nn.L1Loss</code> , which calculates the mean absolute error (MAE) between the input and	1464
1423	* A legend is added to differentiate	the reconstructed output. The <code>reduction='sum'</code> argument	1465
1424	between the training and validation	ensures that the loss is summed	1466
1425	loss curves.	across all elements.	1467
1426	* A title, Loss over training		
1427	epochs, is included to specify the	* The model is set to evaluation mode	1468
1428	purpose of the plot.	using <code>model.eval()</code> to disable	1469
1429	* The plot is displayed using	dropout and batch normalization,	1470
1430	<code>plt.show()</code> .	ensuring consistent predictions.	1471
1431	Purpose of this Visualization: This plot is crucial		1472
1432	for understanding the model's performance during		1473
1433	training:		1474

1480	* The dataset is processed in batches	1527	* The <code>predict</code> function is used to
1481	using a loop. For each batch:	1528	calculate reconstruction losses for
1482	· The true input sequence,	1529	four subsets of data:
1483	<code>seq_true</code> , is passed	1530	· <code>x_train</code> : The training data
1484	through the model to gen-	1531	used to train the autoencoder.
1485	erate the predicted output,	1532	· <code>x_test</code> : The testing data
1486	<code>seq_pred</code> .	1533	used to evaluate the autoen-
1487	· The reconstruction loss	1534	coder's generalization capabil-
1488	between <code>seq_true</code> and	1535	ity.
1489	<code>seq_pred</code> is calculated	1536	· <code>attack_tensor</code> : Data rep-
1490	using the specified criterion.	1537	resenting anomalous samples
1491	· The predicted outputs are ap-	1538	(e.g., attacks) from the dataset.
1492	pended to the <code>predictions</code>	1539	· <code>normal_tensor</code> : Data rep-
1493	list, and the corresponding loss	1540	resenting normal samples from
1494	values are appended to the	1541	the dataset.
1495	<code>losses</code> list.		
1496	* The function returns two outputs:	1542	* For each subset, the <code>predict</code>
1497	· <code>predictions</code> : A list of the	1543	function is called:
1498	predicted outputs generated by	1544	· The first returned value (<code>_</code>)
1499	the model.	1545	contains the predicted outputs,
1500	· <code>losses</code> : A list of the recon-	1546	which are not used here.
1501	struction loss values for each	1547	· The second returned
1502	batch.	1548	value (<code>losses_train</code> ,
1503		1549	<code>losses_test</code> ,
1504	Purpose of the Function: This function is crucial	1550	<code>losses_attack</code> ,
1505	for evaluating the model's reconstruction ability:	1551	<code>losses_normal</code>) stores
1506	* By computing predictions and	1552	the reconstruction losses for
1507	losses, it helps identify how well	1553	each subset.
1508	the model reconstructs input data.		
1509	* Loss values can be analyzed to de-	1554	Purpose of the Code:
1510	tect anomalies, as higher losses of-	1555	* The reconstruction losses provide
1511	ten correspond to unusual or anoma-	1556	a quantitative measure of how well
1512	lous inputs.	1557	the autoencoder performs on each
1513	* It provides a batch-wise processing	1558	subset of data.
1514	mechanism to handle large datasets	1559	* Higher reconstruction losses for
1515	efficiently.	1560	<code>attack_tensor</code> compared to
1516		1561	<code>normal_tensor</code> are expected,
1517	2.9 Loss Calculation for Different Data Sets	1562	as anomalies typically deviate from
1518	This code snippet calculates the reconstruction	1563	the patterns learned by the autoen-
1519	losses for different subsets of data using the trained	1564	coder.
1520	autoencoder model. These losses are essential for	1565	* Losses for <code>x_test</code> indicate the
1521	understanding how well the model reconstructs	1566	model's performance on unseen
1522	data and for identifying anomalies.	1567	data, highlighting its generalization
1523	Description of the Code:	1568	capability.
1524	* <code>x_train.size()</code> : This line re-	1569	* These loss values can be further
1525	trieves the size (dimensions) of the	1570	analyzed to define thresholds for
1526	training dataset, providing insight	1571	anomaly detection or visualized for
	into the number of samples and fea-	1572	better interpretation.
	tures used for training.		

2.10 Visualizing Loss Distributions

This code snippet uses the `seaborn` library to visualize the reconstruction loss distributions for different subsets of data. Visualization is a critical step in understanding the model's performance and in differentiating normal and anomalous data.

Description of the Code:

- * The `sns.kdeplot` function is employed to plot kernel density estimates (KDE) of the reconstruction losses. This technique visualizes the probability density function of the data, providing insights into the distribution of reconstruction losses.

- * Specifically, two distributions are plotted:

- `losses_normal`: The reconstruction losses for normal data are plotted in blue, representing how well the autoencoder reconstructs non-anomalous samples.

- `losses_attack`: The reconstruction losses for anomalous data are plotted in red, highlighting the deviations from normal patterns.

- * Commented-out lines include alternative visualization options, such as histograms (`plt.hist`) and combined KDE and histogram plots (`sns.distplot`), which are not used in this final version.

Purpose of the Code:

- * The KDE plots provide a smooth, intuitive visualization of loss distributions, enabling the comparison of normal and anomalous data.

- * Differences in the distributions help identify potential thresholds for anomaly detection. For instance:

- If the red curve (anomalies) significantly deviates from the blue curve (normal data), this indicates that the autoencoder

effectively distinguishes between normal and anomalous samples.

- Overlapping regions may suggest areas where the model struggles to differentiate anomalies from normal data.

- * These plots can guide the selection of an appropriate loss threshold to classify data as normal or anomalous.

2.11 Brute-Force Optimization of Cutoff Thresholds

This section of the code brute-forces the computation of optimal cutoff thresholds to maximize the accuracy of anomaly detection. The goal is to identify the lower and upper reconstruction loss bounds that best differentiate between normal and anomalous data.

Description of the Code:

- * The code defines two ranges of values:

- `lower_list`: A sequence of candidate lower bounds, generated by adding small increments around an initial lower value (23 in this case).

- `upper_list`: A sequence of candidate upper bounds, similarly generated around an initial upper value (58 in this case).

- * Two empty lists, `ls_1` and `ls_2`, are used to store accuracy scores and their corresponding threshold pairs.

- * For each combination of lower and upper thresholds:

- True positives (TP) are computed as anomalous samples with reconstruction losses within the threshold range (lower to upper).

- False negatives (FN) are anomalous samples outside this range.

1665	· False positives (FP) are normal samples within the threshold range.	* The index of the maximum accuracy score is identified using the <code>np.argmax</code> function applied to the <code>ls_1</code> list, which stores accuracy scores for different threshold pairs.	1710
1666			1711
1667			1712
1668	· True negatives (TN) are normal samples outside this range.		1713
1669			1714
1670			1715
1671	· Accuracy is calculated as the ratio of correctly classified samples (TP + TN) to the total number of samples (TP + TN + FP + FN).	* The optimal lower and upper thresholds corresponding to the highest accuracy are extracted from <code>ls_2</code> .	1716
1672			1717
1673			1718
1674			1719
1675			
1676	* The accuracy scores are appended to <code>ls_1</code> , and their corresponding threshold pairs (<code>lower</code> , <code>upper</code>) are stored in <code>ls_2</code> .	* Using these thresholds:	1720
1677		· True Positives (TP) are computed as anomalous samples with reconstruction losses between the lower and upper thresholds.	1721
1678			1722
1679			1723
1680	Purpose of the Code:		1724
1681			1725
1682	* This brute-force approach explores a wide range of potential threshold pairs to empirically determine the combination that yields the highest accuracy.	· False Negatives (FN) are anomalous samples outside this range.	1726
1683			1727
1684		· False Positives (FP) are normal samples within this range.	1728
1685			1729
1686		· True Negatives (TN) are normal samples outside this range.	1730
1687	* By systematically varying lower and upper bounds, this method ensures thorough coverage of possible threshold combinations, helping to identify optimal cutoff values for classification.		1731
1688			1732
1689		* The True Positive Rate (TPR) is calculated as the ratio of TP to the sum of TP and FN.	1733
1690			1734
1691			1735
1692	Significance:	* The True Negative Rate (TNR) is calculated as the ratio of TN to the sum of TN and FP.	1736
1693			1737
1694	* The identified thresholds can significantly enhance the model's performance in distinguishing between normal and anomalous data, which is critical in anomaly detection systems.	* Accuracy is calculated as the ratio of correctly classified samples (TP + TN) to the total number of samples (TP + TN + FP + FN).	1738
1695			1739
1696			1740
1697			1741
1698			1742
1699			1743
1700	* This approach, while computationally intensive, provides insights into the sensitivity of the model's accuracy to different threshold values.	Purpose of the Code:	1744
1701			1745
1702		* This code evaluates the effectiveness of the optimized thresholds in separating normal and anomalous data based on reconstruction loss.	1746
1703	2.12 Threshold Optimization and Performance Metrics		1747
1704			1748
1705	This section of the code finalizes the threshold optimization process by selecting the optimal thresholds and evaluating the performance of the anomaly detection system.	* By calculating key performance metrics such as TPR, TNR, and accuracy, it provides a comprehensive assessment of the model's ability to detect anomalies.	1749
1706			1750
1707			1751
1708			1752
1709	Description of the Code:	Significance:	1753
			1754

1755	* The metrics calculated in this code	efficient processing. This makes the ap-	1800
1756	(TPR, TNR, accuracy) are essential	proach less accessible for organizations	1801
1757	for understanding the strengths and	with limited computational resources.	1802
1758	weaknesses of the anomaly detec-		
1759	tion system.		
1760	* These metrics highlight the trade-	– Imbalanced Detection Capabilities:	1803
1761	offs between detecting anomalies	The Autoencoder is trained on normal	1804
1762	(sensitivity) and avoiding false	traffic data and relies on reconstruction	1805
1763	alarms (specificity), which are crit-	errors to identify anomalies. This ap-	1806
1764	ical in applications like cybersecu-	proach may lead to higher false positives	1807
1765	rity.	for benign but unusual traffic patterns, as	1808
		well as false negatives for cleverly dis-	1809
		guised malicious activities.	1810
1766	Reference:		
1767	* The precision and recall con-	– Lack of Real-Time Evaluation: The	1811
1768	cepts are further elaborated in the	project has not yet implemented real-	1812
1769	Wikipedia article on Precision and	time monitoring and anomaly detection,	1813
1770	Recall .	which is a critical requirement for many	1814
		practical applications in cybersecurity.	1815
		The current implementation is better	1816
1771	3 Limitations	suited for batch processing and offline	1817
		analysis.	1818
1772	This section outlines the key limitations of the	– Threshold Selection Sensitivity: The	1819
1773	current project and its implementation:	anomaly detection system depends on	1820
1774	– Dependency on Dataset Quality: The	predefined thresholds for reconstruction	1821
1775	performance of the anomaly detection	loss, which can be sensitive to small vari-	1822
1776	system is highly dependent on the qual-	ations in data. Manually or heuristically	1823
1777	ity and representativeness of the dataset.	determining these thresholds might not	1824
1778	While the UNSW-NB15 dataset provides	always result in optimal performance.	1825
1779	diverse attack scenarios, it may not fully		
1780	capture real-world network traffic varia-	– Limited Exploration of Alternative	1826
1781	tions or evolving cyber threats.	Models: While Autoencoders were im-	1827
1782	– Limited Generalizability: The system	plemented, other potentially effective	1828
1783	has been primarily tested on the UNSW-	anomaly detection models, such as Isola-	1829
1784	NB15 dataset, which may limit its ability	tion Forests or ensemble methods, were	1830
1785	to generalize effectively to other datasets	only conceptually explored and not im-	1831
1786	or real-world environments without sig-	plemented or evaluated in this project.	1832
1787	nificant retraining and fine-tuning.		
1788	– Scalability Challenges: The use of Au-	4 Conclusions and Future Work	1833
1789	toencoders and neural networks requires		
1790	substantial computational resources, par-	This project has been a rewarding experience,	1834
1791	ticularly for large-scale networks with	allowing me to explore the application of ma-	1835
1792	high traffic volumes. This may limit the	chine learning techniques in anomaly detec-	1836
1793	scalability of the solution in real-time net-	tion for network security. However, there are a	1837
1794	work environments without optimization	few reflections on what could have been done	1838
1795	or additional hardware.	differently, areas for improvement, and plans	1839
1796	– High Computational Requirements:	for the future.	1840
1797	Training the Autoencoder model is com-	Looking back, starting the project earlier	1841
1798	putationally intensive, requiring access	would have been beneficial. This would have	1842
1799	to GPUs or high-performance CPUs for	allowed me to implement more of the code	1843
		myself and gain a deeper understanding of the	1844
		underlying processes and algorithms. While	1845

1846 I managed to grasp the core concepts and im-
1847 plementation details, having more time to ex-
1848 periment with and fine-tune the models could
1849 have enhanced my learning experience.

1850 There are several ways to improve this project.
1851 One significant enhancement would be to inte-
1852 grate real-time anomaly detection capabilities,
1853 enabling the system to analyze network traffic
1854 as it happens. Additionally, exploring alter-
1855 native machine learning models, such as Iso-
1856 lation Forests or ensemble techniques, could
1857 provide valuable insights and potentially im-
1858 prove the system's performance. Expand-
1859 ing the range of datasets used for evaluation
1860 would also enhance the system's generalizabil-
1861 ity and robustness in real-world scenarios.

1862 I genuinely enjoyed this project. It allowed
1863 me to understand the importance of prepro-
1864 cessing data, designing and training models,
1865 and analyzing their effectiveness. Working
1866 with the UNSW-NB15 dataset and Autoen-
1867 coders helped me develop practical skills in
1868 Python programming and the application of
1869 machine learning frameworks like PyTorch.

1870 Through this project, I learned a great deal
1871 about anomaly detection, network traffic anal-
1872 ysis, and the challenges of developing scalable
1873 machine learning solutions. It deepened my
1874 understanding of the complexities involved in
1875 cybersecurity and the potential of AI-driven
1876 methods to address them.

1877 For future projects in this course, I plan to
1878 focus on Isolation Forests and evaluate their
1879 performance on diverse datasets. Comparing
1880 the results of these models with Autoencoders
1881 could provide a more comprehensive under-
1882 standing of their strengths and weaknesses.
1883 Additionally, incorporating ensemble meth-
1884 ods and advanced feature selection techniques
1885 could make the system even more robust.

1886 Finally, the project was thoroughly enjoyable,
1887 and having more time to experiment and col-
1888 laborate with peers could have made the ex-
1889 perience even better. Overall, this project has
1890 not only been a valuable academic exercise
1891 but also an exciting opportunity to delve into
1892 the practical applications of machine learning
1893 in cybersecurity.

5 References

- 1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
- **GitHub Repository:** Implementation of the Autoencoder for anomaly detection using the UNSW-NB15 dataset. Available at: https://github.com/alik604/cyber-security/blob/master/Intrusion-Detection/UNSW_NB15%20-%20PyTorch%20MLP%20and%20autoEncoder.ipynb.
- **UNSW-NB15 Dataset Description:** Comprehensive details about the UNSW-NB15 dataset for network intrusion detection research. Available at: https://www.unb.ca/cic/datasets/cic-unsw-nb15.html?utm_source=chatgpt.com.
- **UNSW-NB15 Dataset Files:** Training and testing datasets hosted by UNSW. Available at: https://unsw-my.sharepoint.com/personal/z5025758_ad_unsw_edu_au/_layouts/15/onedrive.aspx?id=%2Fpersonal%2Fz5025758%5Fad%5Funsw%5Fedu%5Fau%2FDocuments%2FUNSW%2DNB15%20dataset%2FCSV%20Files%2FTraining%20and%20Testing%20Sets&ga=1.